

Next: [12.4 Discussion and Exercises](#) **Up:** [12. Graphs](#) **Previous:** [12.2 AdjacencyLists: A Graph](#) [Contents](#) [Index](#)

Subsections

- [12.3.1 Breadth-First Search](#)
- [12.3.2 Depth-First Search](#)

12.3 Graph Traversal

In this section we present two algorithms for exploring a graph, starting at one of its vertices, i , and finding all vertices that are reachable from i . Both of these algorithms are best suited to graphs represented using an adjacency list representation. Therefore, when analyzing these algorithms we will assume that the underlying representation is an `AdjacencyLists`.

12.3.1 Breadth-First Search

The breadth-first-search algorithm starts at a vertex i and visits, first the neighbours of i , then the neighbours of the neighbours of i , then the neighbours of the neighbours of the neighbours of i , and so on.

This algorithm is a generalization of the breadth-first traversal algorithm for binary trees (Section [6.1.2](#)), and is very similar; it uses a queue, q , that initially contains only i . It then repeatedly extracts an element from q and adds its neighbours to q , provided that these neighbours have never been in q before. The only major difference between the breadth-first-search algorithm for graphs and the one for trees is that the algorithm for graphs has to ensure that it does not add the same vertex to q more than once. It does this by using an auxiliary boolean array, `seen`, that tracks which vertices have already been discovered.

```
void bfs(Graph g, int r) {
    boolean[] seen = new boolean[g.nVertices()];
    Queue<Integer> q = new SLList<Integer>();
    q.add(r);
    seen[r] = true;
    while (!q.isEmpty()) {
        int i = q.remove();
        for (Integer j : g.outEdges(i)) {
            if (!seen[j]) {
                q.add(j);
                seen[j] = true;
            }
        }
    }
}
```

An example of running `bfs(g,0)` on the graph from Figure 12.1 is shown in Figure 12.4. Different executions are possible, depending on the ordering of the adjacency lists; Figure 12.4 uses the adjacency lists in Figure 12.3.




Figure 12.4: An example of bread-first-search starting at node 0. Nodes are labelled with the order in which they are added to `q`. Edges that result in nodes being added to `q` are drawn in black, other edges are drawn in grey.

Analyzing the running-time of the `bfs(g,i)` routine is fairly straightforward. The use of the `seen` array ensures that no vertex is added to `q` more than once. Adding (and later removing) each vertex from `q` takes constant time per vertex for a total of $O(n)$ time. Since each vertex is processed by the inner loop at most once, each adjacency list is processed at most once, so each edge of `G` is processed at most once. This processing, which is done in the inner loop takes constant time per iteration, for a total of $O(m)$ time. Therefore, the entire algorithm runs in $O(n+m)$ time.

The following theorem summarizes the performance of the `bfs(g,r)` algorithm.

Theorem 12..3 *When given as input a Graph, `g`, that is implemented using the `AdjacencyLists` data structure, the `bfs(g,r)` algorithm runs in $O(n+m)$ time.*

A breadth-first traversal has some very special properties. Calling `bfs(g,r)` will eventually enqueue (and eventually dequeue) every vertex `j` such that there is a directed path from `r` to `j`. Moreover, the vertices at distance 0 from `r` (`r` itself) will enter `q` before the vertices at distance 1, which will enter `q` before the vertices at distance 2, and so on. Thus, the `bfs(g,r)` method visits vertices in increasing order of distance from `r` and vertices that cannot be reached from `r` are never visited at all.

A particularly useful application of the breadth-first-search algorithm is, therefore, in computing shortest paths. To compute the shortest path from r to every other vertex, we use a variant of `bfs(g, r)` that uses an auxiliary array, p , of length n . When a new vertex j is added to q , we set $p[j] = i$. In this way, $p[j]$ becomes the second last node on a shortest path from r to j . Repeating this, by taking $p[p[j]]$, $p[p[p[j]]]$, and so on we can reconstruct the (reversal of) a shortest path from r to j .

12.3.2 Depth-First Search

The depth-first-search algorithm is similar to the standard algorithm for traversing binary trees; it first fully explores one subtree before returning to the current node and then exploring the other subtree. Another way to think of depth-first-search is by saying that it is similar to breadth-first search except that it uses a stack instead of a queue.

During the execution of the depth-first-search algorithm, each vertex, i , is assigned a colour, $c[i]$: **white** if we have never seen the vertex before, **grey** if we are currently visiting that vertex, and **black** if we are done visiting that vertex. The easiest way to think of depth-first-search is as a recursive algorithm. It starts by visiting r . When visiting a vertex i , we first mark i as **grey**. Next, we scan i 's adjacency list and recursively visit any white vertex we find in this list. Finally, we are done processing i , so we colour i black and return.

```
void dfs(Graph g, int r) {
    byte[] c = new byte[g.nVertices()];
    dfs(g, r, c);
}
void dfs(Graph g, int i, byte[] c) {
    c[i] = grey; // currently visiting i
    for (Integer j : g.outEdges(i)) {
        if (c[j] == white) {
            c[j] = grey;
            dfs(g, j, c);
        }
    }
    c[i] = black; // done visiting i
}
```

An example of the execution of this algorithm is shown in Figure [12.5](#).




Figure 12.5: An example of depth-first-search starting at node 0. Nodes are labelled with the order in which they are processed. Edges that result in a recursive call are drawn in black, other edges are drawn in grey.

Although depth-first-search may best be thought of as a recursive algorithm, recursion is not the best way to implement it. Indeed, the code given above will fail for many large graphs by causing a stack overflow. An alternative implementation is to replace the recursion stack with an explicit stack, `s`. The following implementation does just that:

```
void dfs2(Graph g, int r) {
    byte[] c = new byte[g.nVertices()];
    Stack<Integer> s = new Stack<Integer>();
    s.push(r);
    while (!s.isEmpty()) {
        int i = s.pop();
        if (c[i] == white) {
            c[i] = grey;
            for (int j : g.outEdges(i))
                s.push(j);
        }
    }
}
```

In the preceding code, when the next vertex, `i`, is processed, `i` is coloured `grey` and then replaced, on the stack, with its adjacent vertices. During the next iteration, one of these vertices will be visited.

Not surprisingly, the running times of `dfs(g,r)` and `dfs2(g,r)` are the same as that of `bfs(g,r)`:

Theorem 12..4 *When given as input a Graph, `g`, that is implemented using the `AdjacencyLists` data structure, the `dfs(g,r)` and `dfs2(g,r)` algorithms each run in $O(n+m)$ time.*

As with the breadth-first-search algorithm, there is an underlying tree associated with each execution of depth-first-search. When a node `i` goes from `white` to `grey`, this is because `dfs(g,i,c)` was called recursively

while processing some node i' . (In the case of `dfs2(g, r)` algorithm, i is one of the nodes that replaced i' on the stack.) If we think of i' as the parent of i , then we obtain a tree rooted at r . In Figure 12.5, this tree is a path from vertex 0 to vertex 11.

An important property of the depth-first-search algorithm is the following: Suppose that when node i is coloured **grey**, there exists a path from i to some other node j that uses only white vertices. Then j will be coloured first **grey** then **black** before i is coloured **black**. (This can be proven by contradiction, by considering any path P from i to j .)

One application of this property is the detection of cycles. Refer to Figure 12.6. Consider some cycle, C , that can be reached from r . Let i be the first node of C that is coloured **grey**, and let j be the node that precedes i on the cycle C . Then, by the above property, j will be coloured **grey** and the edge (j, i) will be considered by the algorithm while i is still **grey**. Thus, the algorithm can conclude that there is a path, P , from i to j in the depth-first-search tree and the edge (j, i) exists. Therefore, P is also a cycle.




Figure 12.6: The depth-first-search algorithm can be used to detect cycles in G . The node j is coloured **grey** while i is still **grey**. This implies that there is a path, P , from i to j in the depth-first-search tree, and the edge (j, i) implies that P is also a cycle.